

## Last class

## Dependency parsing and logistic regression

Shay Cohen  
(based on slides by Sharon Goldwater)

21 October 2019

Dependency parsing:

- ▶ a fully lexicalized formalism; tree edges connect words in the sentence based on head-dependent relationships.
- ▶ a better fit than constituency grammar for languages with free word order; but has weaknesses (e.g., conjunction).
- ▶ Gaining popularity because of move towards multilingual NLP.

## Today's lecture

- ▶ How do we evaluate dependency parsers?
- ▶ Discriminative versus generative models
- ▶ How do we build a probabilistic model for dependency parsing?

## Example



Parsing **Kim saw Sandy**:

Step	← bot. Stacktop →	Word List	Action	Relations
0	[root]	[Kim,saw,Sandy]	Shift	
1	[root, Kim]	[saw,Sandy]	Shift	
2	[root, Kim, saw]	[Sandy]	LeftArc	Kim←saw
3	[root, saw]	[Sandy]	Shift	
4	[root, saw, Sandy]	[]	RightArc	saw→Sandy
5	[root, saw]	[]	RightArc	root→saw
6	[root]	[]	(done)	

- ▶ Here, top two words on stack are also always adjacent in sentence. Not true in general! (See longer example in JM3.)

## Labelled dependency parsing

- ▶ These parsing actions produce **unlabelled** dependencies (left).
- ▶ For **labelled** dependencies (right), just use more actions: LeftArc(NSUBJ), RightArc(NSUBJ), LeftArc(DOBJ), ...



## Differences to constituency parsing

- ▶ Shift-reduce parser for CFG: not all sequences of actions lead to valid parses. Choose incorrect action → may need to backtrack.
- ▶ Here, all valid action sequences lead to valid parses.
  - ▶ Invalid actions: can't apply LeftArc with root as dependent; can't apply RightArc with root as head unless input is empty.
  - ▶ Other actions may lead to **incorrect** parses, but still **valid**.
- ▶ So, parser doesn't backtrack. Instead, tries to **greedily** predict the correct action at each step.
  - ▶ Therefore, dependency parsers can be very fast (linear time).
  - ▶ But need a good way to predict correct actions (coming up).

## Notions of validity

- ▶ In constituency parsing, valid parse = grammatical parse.
  - ▶ That is, we first define a grammar, then use it for parsing.
- ▶ In dependency parsing, we don't normally define a grammar. Valid parses are those with the properties mentioned earlier:
  - ▶ A single distinguished root word.
  - ▶ All other words have exactly one incoming edge.
  - ▶ A unique path from the root to each other word.

## Summary: Transition-based Parsing

- ▶ **arc-standard** approach is based on simple shift-reduce idea.
- ▶ Can do labelled or unlabelled parsing, but need to train a **classifier** to predict next action, as we'll see.
- ▶ Greedy algorithm means time complexity is linear in sentence length.
- ▶ Only finds **projective** trees (without special extensions)
- ▶ Pioneering system: Nivre's **MALTPARSER**.

## Alternative: Graph-based Parsing

- ▶ Global algorithm: From the fully connected directed graph of all possible edges, choose the best ones that form a tree.
- ▶ **Edge-factored** models: Classifier assigns a nonnegative score to each possible edge; **maximum spanning tree** algorithm finds the spanning tree with highest total score in  $O(n^2)$  time.
- ▶ Pioneering work: McDonald's MSTPARSER
- ▶ Can be formulated as constraint-satisfaction with **integer linear programming** (Martins's TURBOPARSER)
- ▶ Details in JM3, Ch 14.5 (optional).

## Graph-based vs. Transition-based vs. Conversion-based

- ▶ TB: Features in scoring function can look at any part of the stack; no optimality guarantees for search; linear-time; (classically) projective only
- ▶ GB: Features in scoring function limited by factorization; optimal search within that model; quadratic-time; no projectivity constraint
- ▶ CB: In terms of accuracy, sometimes best to first constituency-parse, then convert to dependencies (e.g., STANFORD PARSER). Slower than direct methods.

## Choosing a Parser: Criteria

- ▶ Target representation: constituency or dependency?
- ▶ Efficiency? In practice, both runtime and memory use.
- ▶ Incrementality: parse the whole sentence at once, or obtain partial left-to-right analyses/expectations?
- ▶ Accuracy?

## Probabilistic transition-based dep'y parsing

At each step in parsing we have:

- ▶ Current configuration: consisting of the stack state, input buffer, and dependency relations found so far.
- ▶ Possible actions: e.g., SHIFT, LEFTARC, RIGHTARC.


Probabilistic parser assumes we also have a model that tells us  $P(\text{action}|\text{configuration})$ . Then,

- ▶ Choosing the most probable action at each step (**greedy** parsing) produces a parse in linear time.
- ▶ But it might not be the best one: choices made early could lead to a worse overall parse.

## Recap: parsing as search

Parser is searching through a very large space of possible parses.

- ▶ Greedy parsing is a depth-first strategy.
- ▶ **Beam search** is a limited breadth-first strategy.



## Beam search: basic idea

- ▶ Instead of choosing only the **best** action at each step, choose a few of the best.
- ▶ Extend previous partial parses using these options.
- ▶ At each time step, keep a fixed number of best options, discard anything else.

Advantages:

- ▶ May find a better overall parse than greedy search,
- ▶ While using less time/memory than exhaustive search.

## The agenda

An ordered list of configurations (parser state + parse so far).

- ▶ Items are ordered by score: how good a configuration is it?
- ▶ Implemented using a **priority queue** data structure, which efficiently inserts items into the ordered list.
- ▶ In beam search, we use an agenda with a fixed size (**beam width**). If new high-scoring items are inserted, discard items at the bottom below beam width.

Won't discuss scoring function here; but beam search idea is used across NLP (e.g., in best-first constituency parsing, NNet models.)

## Evaluating dependency parsers

- ▶ How do we know if beam search is helping?
- ▶ As usual, we can evaluate against a gold standard data set. But what evaluation measure to use?

## Evaluating dependency parsers

- ▶ By construction, the number of dependencies is the same as the number of words in the sentence.
- ▶ So we do not need to worry about precision and recall, just plain old accuracy.
- ▶ **Labelled Attachment Score** (LAS): Proportion of words where we predicted the correct head and label.
- ▶ **Unlabelled Attachment Score** (UAS): Proportion of words where we predicted the correct head, regardless of label.

## Building a classifier for next actions

We said:

- ▶ Probabilistic parser assumes we also have a model that tells us  $P(\text{action}|\text{configuration})$ .

Where does that come from?

## Classification for action prediction

We've seen **text classification**:

- ▶ Given (features from) text document, predict the class it belongs to.

Generalized classification task:

- ▶ Given features from observed data, predict one of a set of classes (labels).

Here, **actions** are the labels to predict:


- ▶ Given (features from) the current configuration, predict the next action.

## Training data

Our goal is:

- ▶ Given (features from) the current configuration, predict the next action.

Our corpus contains annotated sentences such as:



Is this sufficient to train a classifier to achieve our goal?

Creating the right training data

- Well, not quite. What we need is a sequence of the correct (configuration, action) pairs.
- ▶ Problem: some sentences may have **more than one** possible sequence that yields the correct parse. (see tutorial exercise)
  - ▶ Solution: JM3 describes rules to convert each annotated sentence to a **unique** sequence of (configuration, action) pairs.<sup>1</sup>

OK, finally! So what kind of model will we train?

<sup>1</sup>This algorithm is called the *training oracle*. An *oracle* is a fortune-teller, and in NLP it refers to an algorithm that always provides the correct answer. Oracles can also be useful for evaluating certain aspects of NLP systems, and we may say a bit more about them later.

Logistic regression

- ▶ Actually, we could use any kind of classifier (Naive Bayes, SVM, neural net...)
- ▶ Logistic regression is a standard approach that illustrates a different type of model: a **discriminative** probabilistic model.
  - ▶ So far, all our models have been **generative**.
- ▶ Even if you have seen it before, the formulation often used in NLP is slightly different from what you might be used to.

Generative probabilistic models

- ▶ Model the joint probability  $P(\vec{x}, \vec{y})$ 
  - ▶  $\vec{x}$ : the observed variables (what we'll see at test time).
  - ▶  $\vec{y}$ : the latent variables (not seen at test time; must predict).

Model	$\vec{x}$	$\vec{y}$
Naive Bayes	features	classes
HMM	words	tags
PCFG	words	tree

Generative models have a “generative story”

- ▶ a probabilistic process that describes how the data were created
  - ▶ Multiplying probabilities of each step gives us  $P(\vec{x}, \vec{y})$ .
- ▶ **Naive Bayes**: For each item  $i$  to be classified, (e.g., document)
  - ▶ Generate its class  $c_i$  (e.g., **SPORT**)
  - ▶ Generate its features  $f_{i1} \dots f_{in}$  conditioned on  $c_i$  (e.g., **ball**, **goal**, **Tuesday**)

## Generative models have a “generative story”

- ▶ a probabilistic process that describes how the data were created
  - ▶ Multiplying probabilities of each step gives us  $P(\vec{x}, \vec{y})$ .
- ▶ **Naive Bayes:** For each item  $i$  to be classified, (e.g., document)
  - ▶ Generate its class  $c_i$  (e.g., *SPORT*)
  - ▶ Generate its features  $f_{i1} \dots f_{in}$  conditioned on  $c_i$  (e.g., *ball, goal, Tuesday*)

Result:

$$P(\vec{c}, \vec{f}) = \prod_i \left[ P(c_i) \prod_j P(f_{ij}|c_i) \right]$$

## Other generative stories

- ▶ **HMM:** For each position  $i$  in sentence,
  - ▶ Generate its tag  $t_i$  conditioned on previous tag  $t_{i-1}$
  - ▶ Generate its word  $w_i$  conditioned on  $t_i$
- ▶ **PCFG:**
  - ▶ Starting from S node, recursively generate children for each phrasal category  $c_i$  conditioned on  $c_i$ , until all unexpanded nodes are pre-terminals (tags).
  - ▶ For each pre-terminal  $t_i$ , generate a word  $w_i$  conditioned on  $t_i$ .

## Inference in generative models

- ▶ At test time, given only  $\vec{x}$ , infer  $\vec{y}$  using Bayes' rule:

$$P(\vec{y}|\vec{x}) = \frac{P(\vec{x}|\vec{y})P(\vec{y})}{P(\vec{x})}$$

- ▶ So, notice we actually model  $P(\vec{x}, \vec{y})$  as  $P(\vec{x}|\vec{y})P(\vec{y})$ .
  - ▶ You can confirm this for each of the previous models.

## Discriminative probabilistic models

- ▶ Model  $P(\vec{y}|\vec{x})$  directly
  - ▶ No model of  $P(\vec{x}, \vec{y})$
  - ▶ No generative story
  - ▶ No Bayes' rule
- ▶ One big advantage: we can use arbitrary features and don't have to make strong independence assumptions.
- ▶ But: unlike generative models, we can't get  $P(\vec{x}) = \sum_{\vec{y}} P(\vec{x}, \vec{y})$ .

## Discriminative models more broadly

- ▶ Trained to **discriminate** right v. wrong value of  $\vec{y}$ , given input  $\vec{x}$ .
- ▶ Need not be probabilistic.
- ▶ Examples: support vector machines, (some) neural networks, decision trees, nearest neighbor methods.
- ▶ Here, we consider only multinomial logistic regression models, which *are* probabilistic.
  - ▶ *multinomial* means more than two possible classes
  - ▶ otherwise (or if lazy) just *logistic regression*
  - ▶ In NLP, also known as **Maximum Entropy** (or **MaxEnt**) models.

## Example task: word sense disambiguation

Remember, logistic regression can be used for any classification task.

The following slides use an example from lexical semantics:

- ▶ Given a word with different meanings (**senses**), can we classify which sense is intended?

I visited the Ford **plant** yesterday.  
The farmers **plant** soybeans in spring.  
This **plant** produced three kilos of berries.

## WSD as example classification task

- ▶ Disambiguate three senses of the target word **plant**
  - ▶  $\vec{x}$  are the words and POS tags in the document the target word occurs in
  - ▶  $y$  is the latent sense. Assume three possibilities:

$y =$	sense
1	Noun: a member of the plant kingdom
2	Verb: to place in the ground
3	Noun: a factory

- ▶ We want to build a model of  $P(y|\vec{x})$ .

## Defining a MaxEnt model: intuition

- ▶ Start by defining a set of **features** that we think are likely to help discriminate the classes. E.g.,
  - ▶ the POS of the target word
  - ▶ the words immediately preceding and following it
  - ▶ other words that occur in the document
- ▶ During training, the model will learn how much each feature contributes to the final decision.



## Defining a MaxEnt model

- Features  $f_i(\vec{x}, y)$  depend on both observed and latent variables. E.g., if **tgt** is the target word:
  - $f_1$ : POS(tgt) = NN &  $y = 1$
  - $f_2$ : POS(tgt) = NN &  $y = 2$
  - $f_3$ : preceding\_word(tgt) = 'chemical' &  $y = 3$
  - $f_4$ : doc\_contains('animal') &  $y = 1$

## Defining a MaxEnt model

- Features  $f_i(\vec{x}, y)$  depend on both observed and latent variables. E.g., if **tgt** is the target word:
  - $f_1$ : POS(tgt) = NN &  $y = 1$
  - $f_2$ : POS(tgt) = NN &  $y = 2$
  - $f_3$ : preceding\_word(tgt) = 'chemical' &  $y = 3$
  - $f_4$ : doc\_contains('animal') &  $y = 1$
- Each feature  $f_i$  has a real-valued weight  $w_i$  (learned in training).
- $P(y|\vec{x})$  is a monotonic function of  $\vec{w} \cdot \vec{f}$  (that is,  $\sum_i w_i f_i(\vec{x}, y)$ ).

## Defining a MaxEnt model

- Features  $f_i(\vec{x}, y)$  depend on both observed and latent variables. E.g., if **tgt** is the target word:
  - $f_1$ : POS(tgt) = NN &  $y = 1$
  - $f_2$ : POS(tgt) = NN &  $y = 2$
  - $f_3$ : preceding\_word(tgt) = 'chemical' &  $y = 3$
  - $f_4$ : doc\_contains('animal') &  $y = 1$
- Each feature  $f_i$  has a real-valued weight  $w_i$  (learned in training).
- $P(y|\vec{x})$  is a monotonic function of  $\vec{w} \cdot \vec{f}$  (that is,  $\sum_i w_i f_i(\vec{x}, y)$ ).
  - To make  $P(y|\vec{x})$  large, we need weights that make  $\vec{w} \cdot \vec{f}$  large.

## Example of features and weights

- Let's look at just two features from the **plant** disambiguation example:
  - $f_1$ : POS(tgt) = NN &  $y = 1$
  - $f_2$ : POS(tgt) = NN &  $y = 2$
- Our classes are:  
{1: member of plant kingdom; 2: put in ground; 3: factory}
- Our example doc ( $\vec{x}$ ):  
[... animal/NN ... chemical/JJ plant/NN ...]

## Two cases to consider

- ▶ Computing  $P(y = 1|\vec{x})$ :
  - ▶ Here,  $f_1 = 1$  and  $f_2 = 0$ .
  - ▶ We would expect the probability to be relatively high.
  - ▶ Can be achieved by having a **positive** value for  $w_1$ .
  - ▶ Since  $f_2 = 0$ , its weight has no effect on the final probability.
- ▶ Computing  $P(y = 2|\vec{x})$ :

## Two cases to consider

- ▶ Computing  $P(y = 1|\vec{x})$ :
  - ▶ Here,  $f_1 = 1$  and  $f_2 = 0$ .
  - ▶ We would expect the probability to be relatively high.
  - ▶ Can be achieved by having a **positive** value for  $w_1$ .
  - ▶ Since  $f_2 = 0$ , its weight has no effect on the final probability.
- ▶ Computing  $P(y = 2|\vec{x})$ :
  - ▶ Here,  $f_1 = 0$  and  $f_2 = 1$ .
  - ▶ We would expect the probability to be close to zero, because sense 2 is a verb sense, and here we have a noun.
  - ▶ Can be achieved by having a large **negative** value for  $w_2$ .
  - ▶ By doing so,  $f_2$  says: "If I am active, do **not** choose sense 2!".

## Classification with MaxEnt

- ▶ Choose the class that has highest probability according to

$$P(y|\vec{x}) = \frac{1}{Z} \exp \left( \sum_i w_i f_i(\vec{x}, y) \right)$$

where

- ▶  $\exp(x) = e^x$  (the monotonic function)
- ▶  $\sum_i w_i f_i$  is the *dot product* of  $\vec{w}$  and  $\vec{f}$ , also written  $\vec{w} \cdot \vec{f}$ .
- ▶ The normalization constant  $Z = \sum_{y'} \exp(\sum_i w_i f_i(\vec{x}, y'))$

## Which features are active?

- ▶ Example doc:  
[... animal/NN ... chemical/JJ plant/NN ...]

$$\begin{array}{lll} P(y = 1|\vec{x}) & \text{will have} & f_1, f_4 = 1 \quad \text{and} \quad f_2, f_3 = 0 \\ P(y = 2|\vec{x}) & & f_2 = 1 \quad \quad \quad f_1, f_3, f_4 = 0 \\ P(y = 3|\vec{x}) & & f_3 = 1 \quad \quad \quad f_1, f_2, f_4 = 0 \end{array}$$

- ▶ Notice that zero-valued features have no effect on the final probability
- ▶ Other features will be multiplied by their weights, summed, then exp.

Feature templates

- ▶ In practice, features are usually defined using **templates**  
    `POS(tgt)=t & y`  
    `preceding_word(tgt)=w & y`  
    `doc_contains(w) & y`
  - ▶ instantiate with all possible POSs *t* or words *w* and classes *y*
  - ▶ usually filter out features occurring very few times
  - ▶ templates can also define real-valued or integer-valued features
- ▶ NLP tasks often have a few templates, but 1000s or 10000s of features

Features for dependency parsing

- ▶ We want the model to tell us  $P(\text{action}|\text{configuration})$ .
- ▶ So *y* is the action, and *x* is the configuration.
- ▶ Features are various combinations of words/tags from stack/input:

Source	Feature templates		
One word	$s_1.w$	$s_1.t$	$s_1.wt$
	$s_2.w$	$s_2.t$	$s_2.wt$
	$b_1.w$	$b_1.w$	$b_0.wt$
Two word	$s_1.w \circ s_2.w$	$s_1.t \circ s_2.t$	$s_1.t \circ b_1.w$
	$s_1.t \circ s_2.wt$	$s_1.w \circ s_2.w \circ s_2.t$	$s_1.w \circ s_1.t \circ s_2.t$
	$s_1.w \circ s_1.t \circ s_2.t$	$s_1.w \circ s_1.t$	

**Figure 14.9** Standard feature templates for training transition-based dependency parsers. In the template specifications  $s_n$  refers to a location on the stack,  $b_n$  refers to a location in the word buffer,  $w$  refers to the wordform of the input, and  $t$  refers to the part of speech of the input.

Summary

- We've discussed
- ▶ Beam search.
  - ▶ Evaluation for probabilistic dependency parsing.
  - ▶ The logistic regression classifier.